Star algorithm for NN ensembling

Sergey Zinchenko
Novosibirsk State University
Russian Federation
zinch.s.e@gmail.com

Dmitry Lishudi
HSE University,
Russian Federation
dlishudi@hse.ru

Abstract

Neural network ensembling is a common and robust way to increase model efficiency. In this paper, we propose a new neural network ensemble algorithm based on Audibert's empirical star algorithm. We provide optimal theoretical minimax bound on the excess squared risk. Additionally, we empirically study this algorithm on regression and classification tasks and compare it to most popular ensembling methods.

1 Introduction

Deep learning has been successfully applied to many types of problems and has reached the state-of-the-art performance. Deep learning models have shown good results in regression analysis and time series forecasting [Qiu+14], computer vision [He+16], as well as in natural language processing [OMK20] and other areas. In many complex problems, such as the Imagenet competition [Den+09], the best results are achieved by ensembles of neural networks, that is, it is often useful to combine the predictions of multiple neural networks to create a new one. The easiest way to ensemble multiple neural networks is to average their predictions [Dru+94]. As shown in work [Kaw16], the number of local minima grows exponentially with the number of parameters. And since modern neural network training methods are based on stochastic optimization, two identical architectures optimized with different initializations will probably converge to different solutions. Such a technique for obtaining neural networks with subsequent construction of an ensemble by majority voting or averaging is used, for example, in article [Car+04].

In addition to the fact that the class of deep neural networks has a huge number of local minima, it is also non-convex. It was shown in work [LM09] that for the procedure of minimizing the empirical risk in a non-convex class of functions, the order of convergence is not optimal. In fact, most modern neural network training methods do just that: they minimize the mean value of some error function on the training set. J.-Y. Audibert proposed the star procedure method, which has optimal rate of convergence of excess squared risk [Aud07]. Motivated by this observation and the huge success of ensembles of neural networks, we propose a modification of the star procedure that will combine the advantages of both methods. In short, the procedure we propose can be described as follows: we run d independent learning processes of neural networks, obtaining empirical risk minimizers $\widehat{g}_1, \ldots, \widehat{g}_d$, freeze their weights, then we initialize a new model and connect all $d + 1$ models with a layer of convex coefficients, after that we start the process of optimizing all non-frozen parameters. This whole procedure can be viewed as a search for an empirical minimizer in all possible d-dimensional simplices spanned by d-minimizers and a class of neural networks. As is known, the minimization of the empirical risk with respect to the convex hull is not optimal in the same way as with respect to the original class of functions. Our method, however, minimizes over some set intermediate between the original class of functions and its convex hull, allowing us to combine the advantages of model ensembling and the star procedure.

One can look at this procedure as a new way to train one large neural network with a block architecture, as well as a new way of aggregating models. In this work, we carry out a theoretical analysis of the

behavior of the proposed algorithm for solving the regression problem with a class of sparse neural networks, and also check the operation of the algorithm in numerical experiments on classification and regression problems.

In addition to this, we take into account that it is impossible to achieve a global minimum in the class of neural networks, and we consider the situation of imprecise minimization.

The *main results* of our work can be formulated as follows:

- A multidimensional modification of the star procedure is proposed.

- We prove that the resulting estimate satisfies the exact oracle inequality. It follows from this estimate that the order of convergence of the algorithm (in terms of sample size n) for a fixed neural network architecture is optimal. Our results improve over the imprecise oracle inequality in [Sch20].

- We give an upper bound on the generalization error for the case of approximate empirical risk minimizers, which implies the stability of our algorithm against minimization errors.

- Based on our algorithm, we propose a new method for training block architecture neural networks, which is quite universal in terms of procedures. We also propose a new way to solve the aggregation problem.

- We illustrate the efficiency of our approach with numerical experiments on real-world datasets.

The rest of this paper is organized as follows.

In Section 2, we make an overview of neural network ensembling methods and briefly discuss the advantages of the star algorithm. In Section 3, following the Schmidt-Hieber notation [Sch20], we define a class of sparse fully connected neural networks and formulate a number of statements from which it follows that the algorithm we proposed has a fast rate of convergence. All proofs are attached in additional materials. In Section 4, we discuss the implementation of our algorithm, point out a number of possible problems, and suggest several modifications to fix them. It also describes the conditions for conducting numerical experiments and presents some of their results. At the end, we offer two possible views on our procedure: a new way to train block neural networks and a fairly flexible model aggregation procedure.

2 Related work

2.1 Ensemble strategies

The main idea of the ensemble is to train several predictors and build a good metamodel on them. There are many techniques for its construction. We present some of them. A more detailed review can be found in the work [Gan+21].

Bagging.

The first of two stages is the generation of several samples with the same distribution as the training one. The next stage is training multiple models and aggregate their predictions. There are cases when the predictions of the constructed models are transferred to another model as new features [Kim+02]. But still, most often, aggregation is performed either by majority voting or by averaging. For example, Kyoungnam Ha et al. showed that the ensemble method for perceptrons based on bagging works better than the individual neural network [HCM05].

Boosting

Another approach to construct ensembles is boosting. The idea is to build one strong model from several weak models by stepwise additive modeling. It was first applied to random trees to construct a so-called random forest. But it has also been applied to deep learning models as well. For example, in the task of recognizing facial expressions [Liu+14], or to improve the predictions of convolutional neural networks [Mog+16].

Snapshots

The main problem in aggregation of deep learning models is the cost of training. Training even one modern model requires a lot of resources, and the ensemble needs a lot. A snapshot technique [Hua+17] and their modification [Gar+18] have been created to combat this problem. In short, during the learning process, the step length in the gradient descent is cyclically changed. This allows a learner to get into various local minima (parameters of which are stored for subsequent aggregation) and, as a consequence, to build an ensemble using a computational budget comparable with the cost of training one model.

Later, Zhang et al. tried combining this idea together with boosting in [Zha+20]. The novelty of this work is that the difference between the models is achieved by re-weighting the objects from the training sample, the step length parameter is changed adaptively, and for aggregation, the best model from the entire training cycle is taken, and not the last one.

Implicit ensembles

In this approach a single model is trained in such a way as to behave like an ensemble. But it requires a much smaller computational budget for training. This is achieved due to the fact that in implicit ensembles the parameters of the models are shared, and their averaging is returned as predictions. For example, an implicit ensemble is the *DropOut* [Sri+14] method or the *DropConnect* [Wan+13] method. During training, each neuron or connection in the neural network has a chance to collapse, and after training, a neural network is returned, the elements of which are weighted by the probabilities of the presence of each element. A similar idea is implemented in the [Hua+16] *Stochastic depth* method for [He+16] residual neural networks. There, the residual blocks are randomly discarded during training, and the transformation goes only through a skip connection.

2.2 Star algorithm

Unlike the ensemble problem, the aggregation problem focuses on building a good predictor in a situation where there are already several ready-made models. The reader is referred to [Nem00; Tsy03] for different types of aggregation. It is important to mention that, in contrast to the two-stage star procedure of Audibert [Aud07], the usual empirical risk minimization procedure among the class of known predictors (or their convex hull) does not necessarily lead to the optimal rate of convergence [LM09]. This result was further developed in [LRS15], where the authors extend the theoretical analysis of the star algorithm to the case of infinite classes using the offset Rademacher's complexity technique. It was also shown in the [Vij21] that these results can be generalized to other loss functions. In particular, this means that the star procedure can be applied to more than just regression problems.

3 Theoretical results

In this chapter, we will formulate our theoretical results for solving the regression problem with a class of sparse fully connected neural networks using our proposed $Star_d$ procedure.

We have a $S_n = (\mathbf{X}_i, Y_i)_{i=1}^n$ sample of i.i.d. input-output pairs $(\mathbf{X}_i, Y_i) \in \mathcal{X} \times \mathcal{Y}$ distributed according to some unknown distribution $\mathcal{P}$. We also chose a certain family of solutions $\mathcal{F}$. Our goal is to build a new predictor $\widehat{f}$ minimizing the excess risk

$$\mathcal{E}(\widehat{g}) := \mathbb{E}(\widehat{g} - Y)^2 - \inf_{f \in \mathcal{F}} \mathbb{E}(f - Y)^2.$$

Let $\widehat{\mathbb{E}}$ denote the empirical expectation operator

$$\widehat{\mathbb{E}}(f) := \frac{1}{n} \sum_{i=1}^n f(\mathbf{X}_i)$$

and call $\widehat{g} \in \mathcal{F}$ a Δ-*empirical risk minimizer in* $\mathcal{F}$ if the following inequality holds

$$\widehat{\mathbb{E}}(\widehat{g} - Y)^2 \le \min_{f \in \mathcal{F}} \widehat{\mathbb{E}}(f - Y)^2 + \Delta.$$

We suggest the next two step procedure. In the first, calculate $\{\widehat{g}_i\}_{i=1}^d$ – different Δ_1-empirical risk minimizers in $\mathcal{F}$. And then look for a Δ_2-empirical risk minizer in the next set:

$$Star_d(\mathcal{F}, \widehat{g}_1, \ldots \widehat{g}_d) := \left\{ \sum_{i=1}^d \lambda_i \widehat{g}_i + \underbrace{\left(1 - \sum_{i=1}^d \lambda_i\right) f}_{\lambda} \,\middle|\, \lambda_i, \lambda \in [0,1]; \; f \in \mathcal{F} \right\}. \tag{1}$$

We will call the found function $\widehat{f} = \widehat{f}(\mathcal{F}, d, \Delta_1, \Delta_2)$ as $Star_d$ *estimator*. The main result of our work is the proof that the proposed estimator has an optimal excess risk convergence rate in the case when $\mathcal{F}$ is a class of fully connected neural networks. It is defined by the choice of the activation function $\sigma : \mathbb{R} \to \mathbb{R}$ and the network architecture. We study neural network with activation function ReLu:

$$\sigma(x) := max(x, 0).$$

For $\mathbf{v} = (v_1, \ldots, v_r) \in \mathbb{R}^r$ define shifted activation function $\sigma_{\mathbf{v}} : \mathbb{R}^r \to \mathbb{R}^r$:

$$\sigma_{\mathbf{v}}(\mathbf{x}) := \left(\sigma(x_i - v_i) \right)_{i=1}^r.$$

The network architecture $(L, \mathbf{p})$ consists of a positive integer L called the *number of hidden layers or depth* and a *width vector* $\mathbf{p} = (p_0, \ldots, p_{L+1}) \in \mathbb{N}^{L+2}$. A neural network with network architecture $(L, \mathbf{p})$ is then any function of the form

$$f(\mathbf{x}) = W_L \sigma_{\mathbf{v}_L} W_{L-1} \sigma_{\mathbf{v}_{L-1}} \ldots W_1 \sigma_{\mathbf{v}_1} W_0 \mathbf{x}, \tag{2}$$

where W_j is a $p_{j+1} \times p_j$ matrix and $\mathbf{v}_i \in \mathbb{R}^{p_i}$ is a shift vector.

We will focus on the case when the model parameters satisfy some constraint. Denote $\|W_j\|_\infty$ the maximum-entry norm of W_j, $\|W_j\|_0$ the number of non-zero/active entries of W_j then the space of network functions with given network architecture and network parameters bounded by one is

$$\mathcal{F}(L, \mathbf{p}) := \left\{ f \text{of the form (2)} : \max_{j=0,\ldots,L} \|W_j\|_\infty \vee \|\mathbf{v}_j\|_\infty \leq 1 \right\}$$

and the s-sparse networks are given by

$$\mathcal{F}(L, \mathbf{p}, s) := \left\{ f \in \mathcal{F}(L, \mathbf{p}) : \sum_{j=0}^L \|W_j\|_0 + \|\mathbf{v}_j\|_0 \leq s \right\}.$$

Let's denote by $\mathcal{N}_\infty(\mathcal{F}, \varepsilon)$, $\mathcal{N}_2(\mathcal{F}, \varepsilon)$ the size of the ε-net of $\mathcal{F}$ in the metric space L_∞ and L_2, respectively. Then from Lemma 5 in [Sch20] we have

$$\forall f \in \mathcal{F}(L, \mathbf{p}, s): \|f\|_\infty \leq V(L+1) \tag{3}$$

and

$$\log \mathcal{N}_2(\mathcal{F}(L, \mathbf{p}, s), \delta) \leq \log \mathcal{N}_\infty(\mathcal{F}(L, \mathbf{p}, s), \delta) \leq (s+1) \log(2\delta^{-1}(L+1)V^2), \tag{4}$$

where

$$V := \prod_{l=0}^{L+1} (p_l + 1). \tag{5}$$

Also let define the risk-minimizer in $\mathcal{F}$ and some sets:

$$Hull_d(\mathcal{F}) := \left\{ \sum_{i=1}^d \lambda_i (f_i - f) \,\middle|\, \lambda_i \in [0,1]; \; \sum_{i=1}^d \lambda_i \leq 1; \; f, f_1 \ldots f_d \in \mathcal{F} \right\}, \tag{6}$$

$$f^* := \arg\min_{f \in \mathcal{F}} \mathbb{E}(f(\mathbf{X}) - Y)^2, \quad \mathcal{H} := \mathcal{F} - f^* + Hull_d(\mathcal{F}). \tag{7}$$

Notice, that $Star_d$ estimator $\widehat{f}$ lies in $\mathcal{H} + f^*$. With the introduced notation, one of our main results is stated as follows.

4

Theorem 3.1. *Let $\hat{f}$ be a $Star_d$ estimator and $\mathcal{H}$ be the set defined in 7 for $\mathcal{F} = \mathcal{F}(L, \boldsymbol{p}, s)$. The following expectation bound on excess loss holds:*

$$\mathbb{E}\,\mathcal{E}(\hat{f}) \leq 2(F' + V(L+1)) \cdot \left[\frac{K}{n} + M \cdot \frac{c_{A.8}d\,s\log\left(VLnd\right)}{n}\right] + 4[\Delta_1 + \Delta_2],$$

where

$$C = \min\left\{\frac{c_{A.2}}{4F'}, \frac{c_{A.2}}{4V(L+1)(2+c_{A.2})}\right\}, \quad F' = \sup_{\mathcal{F}}|Y - f|_\infty,$$

$$K := 2\left(\sqrt{\sum_{i=1}^{n}\xi^2/n} + C\right), \quad M := \sup_{h\in\mathcal{H}\setminus\{0\}}4\frac{\sum_{i=1}^{n}h(\boldsymbol{X}_i)^2\xi_i^2}{C\sum_{i=1}^{n}h(\boldsymbol{X}_i)^2}. \tag{8}$$

Thus, the expectation of excess risk for $Star_d$ estimator is limited for a fixed architecture neural network to a value of the order

$$\mathcal{O}\left(\frac{s\,d\,\log(V\,L\,n)}{n} + \Delta_1 + \Delta_2\right).$$

In order to formulate an upper bound for the excess risk, performed with a high probability, we need to impose some constraints on the class of functions.

Definition 3.2 (Lower Isometry Bound). *Class $\mathcal{F}$ satisfies the lower isometry bound with some parameters $0 < \eta < 1$ and $0 < \delta < 1$ if*

$$\mathbb{P}\left(\inf_{f\in\mathcal{F}\setminus\{0\}}\frac{1}{n}\sum_{i=1}^{n}\frac{f^2(\boldsymbol{X}_i)}{\mathbb{E}\,f^2} \geq 1 - \eta\right) \geq 1 - \delta$$

for all $n \geq n_0(\mathcal{F}, \delta, \eta)$, where $n_0(\mathcal{F}, \delta, \eta)$ depends on the complexity of the class.

Now we are ready to formulate the main theoretical result of the paper.

Theorem 3.3. *Let $\hat{f}$ be a $Star_d$ estimator and let $\mathcal{H}$ be the set defined in 7 for $\mathcal{F} = \mathcal{F}(L, \boldsymbol{p}, s)$. Assume for $\mathcal{H}$ the lower isometry bound in Definition 3.2 holds with $\eta_{lib} = c_{A.2}/4$ and some $\delta_{lib} < 1$. Let $\xi_i = Y_i - f^*(\boldsymbol{X}_i)$. Define*

$$A := \sup_{h\in\mathcal{H}}\frac{\mathbb{E}\,h^4}{(\mathbb{E}\,h^2)^2} \quad \text{and} \quad B := \sup_{\boldsymbol{X},Y}\mathbb{E}\,\xi^4.$$

Then there exist 3 absolute constants $c'_{A.10}, c_{\tilde{A}.10}, c_{A.10} > 0$ (which only depend on $c_{A.2}$), such that

$$\mathbb{P}\left(\mathcal{E}(\hat{f}) > 4(D + \Delta_1 + \Delta_2)\right) \leq 4(\delta_{lib} + \delta)$$

as long as $n > \frac{16(1-c'_{A.10})^2 A}{c'^2_{A.10}} \vee n_0(\mathcal{H}, \delta_{lib}, c_{A.10}/4)$, where

$$K := \left(\sqrt{\sum_{i=1}^{n}\xi^2/n} + 2c_{\tilde{A}.10}\right), \quad M := \sup_{h\in\mathcal{H}\setminus\{0\}}\frac{\sum_{i=1}^{n}h(\boldsymbol{X}_i)^2\xi_i^2}{c_{\tilde{A}.10}\sum_{i=1}^{n}h(\boldsymbol{X}_i)^2},$$

$$D := \max\left(\frac{K}{n} + M \cdot \frac{c_{A.8}d\,s\log\left(VLnd\right) + \log\frac{1}{\delta}}{n}, \frac{32\sqrt{AB}}{c'_{A.10}}\frac{1}{n}\right)$$

and $c_{A.8}$ is an independent constant.

That is, with some assumptions on the class of neural networks $\mathcal{F}$, we again obtained the order $\mathcal{O}\left(\frac{\log n\mathcal{N}}{n}\right)$ of convergence of the excess risk. Note that in the general case for an infinite class functions such an asymptotic rate with respect to the sample size n is *unimprovable* (for example, see [RST17, Theorem 3 and Section 5]).

4 Experiments

4.1 Realization

The proposed $Star_d$ procedure can be represented by the following pseudocode (see Algorithm 1). Since in practice we will not be able to give an infinite number of predictors, in fact the algorithm takes the architecture of the neural network as $\mathcal{F}$, without fixing its parameters. The $calculate\_erm$ procedure is some optimization process that reduces the empirical risk. Also, in practice, it is impossible to search for a

---
Algorithm 1: $Star_d$ algorithm

Data: class of functions $\mathcal{F}$ and parameter d
Result: estimator $\widehat{f}$
$\widehat{g}_1 \leftarrow calculate\_erm(set = \mathcal{F}, seed = seed_1)$;
$\ldots$
$\widehat{g}_d \leftarrow calculate\_erm(set = \mathcal{F}, seed = seed_d)$;
$\widehat{f} \leftarrow calculate\_erm(set = Star_d, seed = seed)$;
return $\widehat{f}$

---

global empirical risk minimizer in the space of neural networks, which is why we introduced the concept of Δ-minimizers. As follows from our results, the more accurate the optimization is at each step of the algorithm, the higher the accuracy guarantee of the final predictor $\widehat{f}$.

Despite the fact that the second step of the star algorithm requires an optimization procedure over some complex set $Star_d$, this is fairly easy to implement in practice (see Figure 1). Suppose that we have fixed some architecture of estimator (black block), then in the first step we independently optimize the weights of the blocks $\widehat{g}_i$, freeze them and in the second step we add a new block f, connecting all of them by convex layer λ_i (red elements) and optimize them. This actually iterates over all possible simplices, optimizing the weights of the lower block, and all possible points within the simplex, optimizing the convex weights. To ensure the convexity coefficients in practice we took softmax from the parameter vector.


Figure 1: Illustrate $Star_d$ algorithm on NN

4.2 Competitors and other details

Looking at the structure of the proposed star procedure (let us call it **Classic Star (no warm-up)**), a desire naturally arises to compare it with the following models. First, with the model where all $d + 1$ blocks together with the connecting *linear* layer are trained simultaneously (**Big NN**), and second, where all $d + 1$ blocks are trained independently, and then their predictions are averaged (**Ensemble**).

But the star procedure has a number of possible problems. Training even one neural network is a rather complicated process, while in our algorithm it is required to train $d + 1$ predictors. To solve this problem, we propose to train the $d + 1$ models in the ensemble method and the first d models in the Classic Star procedure using the snapshot technique, calling such approaches **Snap Ensemble** and **Snap Star**, respectively.

Additionally, at the last step of the star procedure, when optimizing the weights of the last model and convex coefficients, the minimizer has the opportunity to set the contribution of the last model to zero and focus on the aggregation of already trained d models. We offer variations of the star algorithm that try to deal with these. It is possible to train the last snapshot (model d) from the Snap Star as a $d + 1$ model, this is what we call **Snap Star (shot warm-up)**, also can add a new model to the Snap Star or Classic Star and spend part of the computational budget on training only its weights, the so-called warm-up, and get **Snap Star (new warm-up)** and **Classic Star (new warm-up)**, respectively. We compare all these methods with each other on a number of tasks. The amount of the computational budget for training a separate block we will denote as $epochs$, in the snapshot technique *learning rate* decreases with cyclical cosine annealing [LH16]. Architecture of neural networks and hardware for training are be reported for each task separately. We ran all experiments using the PyTorch[1] library. For the purposes of reproducing the results, the code and extended tables with results are publicly available at repository[2].

---
[1] https://pytorch.org/library
[2] https://github.com/mordiggian174/star-ensembling

4.3 Datasets

Boston House Pricing We conduct the first experiment on the Boston House Pricing dataset [HR78], whose task is to predict the value of real estate according to some characteristics. The ratio of training and test samples is equal to $7 : 3$. Standard scaler was used as preprocessing, batch size is 32. A small fully connected ReLu neural network of 4 layers was chosen as the architecture of the neural network, the number of neurons on the first layer is 128, then with the growth of the layer it decreases by 2 times, DropOut with parameter p and batch normalization are applied between the layers. For warming up a new model, 40% of the budget for block training is spent. Training was done using Adam [KB14] on the CPU Apple M1. Error results on the test set were averaged over 5 runs. Some of the data obtained in the course of numerical experiments we present in the Table 1. Hereinafter, we highlight in **bold** the best results with a fixed parameter of the number of models d.

It is worth adding that without the use of the DropOut technique, the usual Ensemble with some parameters proved to be quite good (e.g. **10.236** MSE after 100 $epochs$ with $d = 5$, $p = 0.0$, $lr = 0.01$). In total, about 100 different sets of parameters were tried on this launch, we attach part of the results in Appendix B, and the full versions can be found in the repository.

Name	d	MSE	MAE	R^2	TRAIN MSE	TIME (sec)
Snap Star (shot warm-up)	5	11.155±0.967	2.185	0.866	1.232	52.4
Snap Star (new warm-up)	5	10.677±1.034	2.148	0.872	1.306	49.2
Snap Ensemble	5	11.123±0.615	2.206	0.866	1.728	43.4
Ensemble	5	10.172±0.621	2.164	0.878	2.133	43.2
Classic Star (no warm-up)	5	**9.650±0.911**	2.092	**0.884**	1.373	52.0
Classic Star (new warm-up)	5	9.664±0.324	**2.086**	**0.884**	1.374	45.8
Big NN	5	12.117±1.601	2.326	0.854	2.160	28.4
Snap Star (shot warm-up)	4	10.745±0.640	2.154	0.871	1.192	43.0
Snap Star (new warm-up)	4	10.862±0.537	2.162	0.870	1.401	43.0
Snap Ensemble	4	11.169±0.558	2.205	0.866	1.679	36.0
Ensemble	4	9.967±0.244	2.144	0.880	1.833	35.8
Classic Star (no warm-up)	4	9.887±0.501	**2.111**	0.881	1.326	43.6
Classic Star (new warm-up)	4	**9.741±0.790**	2.114	**0.883**	1.400	38.2
Big NN	4	11.856±0.429	2.323	0.858	1.937	24.0

Table 1: BOSTON HOUSE PRICING. Part of results at 200 epochs, $p = 0.1$, $lr = 0.01$

Fashion Mnist The second experiment was carried out on the Fashion Mnist dataset [XRV17], which consists of $70,000$ images (28×28 pixels). It is required to classify images by clothing classes. The ratio of training and test samples is equal to $6 : 1$. No scaler is used, batch size is 64. A simple convolutional network LeNet was chosen as a solution to this task. For warming up a new model, 40% of $epochs$ is again spent. Training was done using Adam on the CPU Apple M1. Error results on the test set were averaged over 3 runs.

As can be seen from the results (see Table 2), the algorithm proposed by us is again in the lead. We also did one relatively hard run with parameters $d = 5$ and 25 $epochs$ for train, the Classic Star (new warm-up) won with 92.2% accuracy. At the moment, they take $10 - 11th$ place in the leaderboard[3] for this dataset. For comparison, the model from the $10th$ place has more than $380K$ parameters. It is worth noting that we used fairly simple models (about $266K$ parameters) and training took about 40 minutes on a regular laptop. The results of training by our method also exceed the indicators Resnet18[4] with over $2M$ parameters and 92% accuracy.

Million Song For the last experiment we use UCI [DG17] subset of Million Song dataset[Ber+11], containing 515345 songs, with 90 features each. Dataset's features are timbre average and covariance for every song, target is the year the song was released. Dataset is divided intro train and test in proportion $9 : 1$, where test is the last 10% of the dataset. Features and targets were standard normalized prior to model training, batch size is 128. Base model for this task is fully-connected 4-layer ReLu neural network with layer dimensions $90, 120, 20, 1$. After each hidden layer, batch

---

[3]https://paperswithcode.com/sota/image-classification-on-fashion-mnist
[4]https://github.com/kefth/fashion-mnist

Name	d	accuracy	entropy	TIME (sec)
Snap Star (shot warm-up)	3	0.902±0.001	0.354±0.017	1021.3
Snap Star (new warm-up)	3	0.902±0.001	0.348±0.015	940.3
Snap Ensemble	3	0.904±0.002	0.274±0.006	819.7
Ensemble	3	0.910±0.003	0.247±0.004	818.7
Classic Star (no warm-up)	3	**0.913±0.002**	0.245±0.007	1014.7
Classic Star (new warm-up)	3	**0.913±0.004**	**0.243±0.008**	855.3
Big NN	3	0.905±0.004	0.297±0.011	643.0
Snap Star (shot warm-up)	2	0.906±0.004	0.292±0.015	751.0
Snap Star (new warm-up)	2	0.907±0.005	0.284±0.012	697.0
Snap Ensemble	2	0.908±0.004	0.266±0.011	616.0
Ensemble	2	**0.912±0.004**	0.251±0.004	614.3
Classic Star (no warm-up)	2	0.910±0.000	0.245±0.003	747.3
Classic Star (new warm-up)	2	0.911±0.002	**0.244±0.004**	613.3
Big NN	2	0.902±0.001	0.303±0.007	498.3

Table 2: FASHION MNIST. Part of results at 15 epochs, $lr = 0.001$

Name	d	MSE	MAE	R^2	Train MSE	TIME (sec)
Snap Star (shot warm-up)	5	76.70 ± 0.28	5.97	0.358	68.05	2369
Snap Star (new warm-up)	5	76.16 ± 0.15	5.96	0.363	69.39	2076
Snap Ensemble	5	76.48 ± 0.18	6.01	0.360	68.15	1766
Ensemble	5	75.74 ± 0.06	5.99	0.366	68.99	1955
Classic Star (no warm-up)	5	75.97 ± 0.20	6.02	0.365	70.88	2150
Classic Star (new warm-up)	5	$\mathbf{75.62 \pm 0.06}$	5.96	0.367	70.29	2176
Big NN	5	76.23 ± 0.19	5.95	0.362	71.70	1166
Snap Star (shot warm-up)	4	76.74 ± 0.20	5.96	0.358	68.44	2026
Snap Star (new warm-up)	4	76.22 ± 0.18	5.97	0.362	69.74	1753
Snap Ensemble	4	76.48 ± 0.19	6.01	0.360	68.41	1473
Ensemble	4	75.78 ± 0.07	5.99	0.366	69.01	1636
Classic Star (no warm-up)	4	76.02 ± 0.09	6.01	0.364	71.11	1811
Classic Star (new warm-up)	4	$\mathbf{75.76 \pm 0.17}$	5.96	0.366	70.41	1823
Big NN	4	76.29 ± 0.29	5.96	0.362	71.74	980

Table 3: MILLIION SONG. Part of results at 30 epochs

normalization and DropOut with probability of 0.2 are applied. Snapshots used SGD optimizer starting at *learning rate* of 0.1. All other models used Adam optimizer with *learning rate* of 0.001. In contrast to previous experiments, for warming up a new model, 10% of $epochs$ is spent and shot warm-up used next fully trained snapshot instead of a last shot (that is, the total training budget is $(d + 2) \cdot epochs$). Training was done on the NVIDIA GeForce GTX 1060 GPU (6gb) / Intel Core i5-7500. Error results were averaged over 5 runs. Some results are shown in the Table 3.

4.4 General patterns

From the obtained results it follows that the best results are achieved by the proposed algorithm with the warm-up of the last model. We believe that the reason for this is that, with a weakly trained last model, it is easier for the algorithm to nullify the contribution of the last model and focus on the predictions of well-trained d estimators. Although results of star algorithm with snapshot technique are mediocre for high number of $epochs$, it shows quite good results for smaller number of $epochs$ (see Tables with a small number of $epochs$ in Appendix B).

5 Discussion and limitations

The proposed algorithm performs well in the classification problem with cross-entropy loss, although this paper only presents a theoretical analysis of regression with squared loss. The practical implementation of the algorithm can be sensitive to the initialization of the last block, as shown by the

warm-up method. And experiments on larger models were not carried out due to high computational costs. Some solutions to these problems are offered in Section 4.2.

In fact, the star estimator we proposed is a multidimensional analogue of the Audibert's algorithm. It combines optimal orders as a solution to the aggregation problem of model selection, and at the same time behaves like an ensemble method. This decision can be viewed from two sides at once.

$Star_d$ algorithm as a new learning algorithm for neural networks of block architecture

It is worth noting that if we spend a fixed amount of computing resources B for each call to the optimization process $calculate\_erm$, then the total budget of our algorithm is about $(d+1) \cdot B$. But the surprising fact is that the result obtained is able to compete with other methods for training the final large neural network from $d+1$ blocks, although our theoretical analysis guarantees optimality only in comparison with the best single block architecture model. Thus, the procedure we proposed can be perceived as a new method for training neural networks with block architecture.

$Star_d$ algorithm as a new way of model aggregation

Also note that the predictors $\widehat{g}_i$ need not be trained in the first step. Then the $Star_d$ algorithm can be perceived as an algorithm for aggregating these models. It will consist of the following: a new predictive model f is added, a connecting layer, and the process of optimization by a parameter is started. At the same time, generally speaking, it is not necessary to have all blocks be of the same architecture. As intuition suggests, the main thing is that the expressive abilities of those classes of solutions to which the predictors given to us will relate should be approximately equal. Then it will be possible to formally consider the union of those decision classes to which each of the predictors belongs, and consider them as Δ-minimizers from the following class $\mathcal{F} = \bigcup_i \mathcal{F}_i$, where given predictors $\widehat{g}_i \in \mathcal{F}_i$ (which may be heterogeneous).

6 Acknowledgments

References

[HR78]   David Harrison Jr and Daniel L Rubinfeld. "Hedonic housing prices and the demand for clean air". In: *Journal of environmental economics and management* 5.1 (1978), pp. 81–102.

[Dru+94] Harris Drucker et al. "Boosting and other ensemble methods". In: *Neural Computation* 6.6 (1994), pp. 1289–1301.

[Nem00]   Arkadi Nemirovski. "Topics in non-parametric statistics". In: *Lectures on probability theory and statistics (Saint-Flour, 1998)* 1738 (2000), pp. 85–277.

[Kim+02] Hyun-Chul Kim et al. "Support vector machine ensemble with bagging". In: *International workshop on support vector machines*. Springer. 2002, pp. 397–408.

[Tsy03]   Alexandre B. Tsybakov. "Optimal Rates of Aggregation". In: *COLT*. 2003.

[Car+04] Rich Caruana et al. "Ensemble selection from libraries of models". In: *Proceedings of the twenty-first international conference on Machine learning*. 2004, p. 18.

[HCM05]   Kyoungnam Ha, Sungzoon Cho, and Douglas MacLachlan. "Response models based on bagging neural networks". In: *Journal of Interactive Marketing* 19.1 (2005), pp. 17–30.

[Aud07]   J.-Y. Audibert. "Progressive mixture rules are deviation suboptimal". In: *NeurIPS* (2007).

[Den+09] Jia Deng et al. "Imagenet: A large-scale hierarchical image database". In: *2009 IEEE conference on computer vision and pattern recognition*. Ieee. 2009, pp. 248–255.

[LM09]   Guillaume Lecué and Shahar Mendelson. "Aggregation via empirical risk minimization". In: *Probability theory and related fields* 145.3 (2009), pp. 591–613.

[Ber+11] Thierry Bertin-Mahieux et al. "The Million Song Dataset". In: *Proceedings of the 12th International Conference on Music Information Retrieval (ISMIR 2011)*. 2011.

[Wan+13] Li Wan et al. "Regularization of neural networks using dropconnect". In: *International conference on machine learning*. PMLR. 2013, pp. 1058–1066.

[KB14]   Diederik P Kingma and Jimmy Ba. "Adam: A method for stochastic optimization". In: *arXiv preprint arXiv:1412.6980* (2014).

[Liu+14]   Ping Liu et al. "Facial expression recognition via a boosted deep belief network". In: *Proceedings of the IEEE conference on computer vision and pattern recognition*. 2014, pp. 1805–1812.

[Qiu+14]   Xueheng Qiu et al. "Ensemble deep learning for regression and time series forecasting". In: *2014 IEEE symposium on computational intelligence in ensemble learning (CIEL)*. IEEE. 2014, pp. 1–6.

[Sri+14]   Nitish Srivastava et al. "Dropout: A Simple Way to Prevent Neural Networks from Overfitting". In: *Journal of Machine Learning Research* 15.56 (2014), pp. 1929–1958. URL: http://jmlr.org/papers/v15/srivastava14a.html.

[LRS15]   Tengyuan Liang, Alexander Rakhlin, and Karthik Sridharan. "Learning with square loss: Localization through offset rademacher complexity". In: *Conference on Learning Theory*. PMLR. 2015, pp. 1260–1285.

[He+16]   Kaiming He et al. "Deep residual learning for image recognition". In: *Proceedings of the IEEE conference on computer vision and pattern recognition*. 2016, pp. 770–778.

[Hua+16]   Gao Huang et al. "Deep networks with stochastic depth". In: *European conference on computer vision*. Springer. 2016, pp. 646–661.

[Kaw16]   Kenji Kawaguchi. "Deep learning without poor local minima". In: *Advances in neural information processing systems* 29 (2016).

[LH16]   Ilya Loshchilov and Frank Hutter. "Sgdr: Stochastic gradient descent with warm restarts". In: *arXiv preprint arXiv:1608.03983* (2016).

[Mog+16]   Mohammad Moghimi et al. "Boosted convolutional neural networks." In: *BMVC*. Vol. 5. 2016, p. 6.

[DG17]   Dheeru Dua and Casey Graff. *UCI Machine Learning Repository*. 2017. URL: http://archive.ics.uci.edu/ml.

[Hua+17]   Gao Huang et al. *Snapshot Ensembles: Train 1, get M for free*. 2017. DOI: 10.48550/ARXIV.1704.00109. URL: https://arxiv.org/abs/1704.00109.

[RST17]   Alexander Rakhlin, Karthik Sridharan, and Alexandre B Tsybakov. "Empirical entropy, minimax regret and minimax risk". In: *Bernoulli* 23.2 (2017), pp. 789–824.

[XRV17]   Han Xiao, Kashif Rasul, and Roland Vollgraf. "Fashion-mnist: a novel image dataset for benchmarking machine learning algorithms". In: *arXiv preprint arXiv:1708.07747* (2017).

[Gar+18]   Timur Garipov et al. *Loss Surfaces, Mode Connectivity, and Fast Ensembling of DNNs*. 2018. DOI: 10.48550/ARXIV.1802.10026. URL: https://arxiv.org/abs/1802.10026.

[OMK20]   Daniel W Otter, Julian R Medina, and Jugal K Kalita. "A survey of the usages of deep learning for natural language processing". In: *IEEE transactions on neural networks and learning systems* 32.2 (2020), pp. 604–624.

[Sch20]   Johannes Schmidt-Hieber. "Nonparametric regression using deep neural networks with ReLU activation function". In: *The Annals of Statistics* 48.4 (2020), pp. 1875–1897.

[Zha+20]   Wentao Zhang et al. "Snapshot boosting: a fast ensemble framework for deep neural networks". In: *Science China Information Sciences* 63.1 (2020), pp. 1–12.

[Gan+21]   M. A. Ganaie et al. *Ensemble deep learning: A review*. 2021. DOI: 10.48550/ARXIV.2104.02395. URL: https://arxiv.org/abs/2104.02395.

[Vij21]   Suhas Vijaykumar. "Localization, Convexity, and Star Aggregation". In: *Advances in Neural Information Processing Systems* 34 (2021).

A   Proofs

The combination of the following 2 Lemmas is a generalization of the geometric inequality proved by Liang et al. [LRS15]. In many respects the scheme of the proof is similar.

Lemma A.1. *(Geometric inequality for the exact $Star_d$ estimator in the second step)*
Let $\hat{g}_1 \ldots \hat{g}_d$ be Δ_1-empirical risk minimizers from the first step of the $Star_d$ procedure, $\widetilde{f}$ be the exact minimizer from the second step of the $Star_d$ procedure. Then, for $c_{A.1} = \frac{1}{18}$ the following inequality holds:

$$\widehat{\mathbb{E}}(h - Y)^2 - \widehat{\mathbb{E}}(\widetilde{f} - Y)^2 \geq c_{A.1}\,\widehat{\mathbb{E}}(\widetilde{f} - h)^2 - 2\Delta_1. \tag{9}$$

Proof. For any function f, g we denote the empirical ℓ_2 distance to be $\|f\|_n := \left[\widehat{\mathbb{E}}\,f^2\right]^{\frac{1}{2}}$, empirical product to be $\langle f, g\rangle_n := \widehat{\mathbb{E}}\,[fg]$ and the square of the empirical distance between $\mathcal{F}$ and Y as r_1. By definition of $Star_d$ estimator for some $\lambda \in [0; 1]$ we have:

$$\widetilde{f} = (1 - \lambda)\widehat{g} + \lambda f,$$

where $\widehat{g}$ lies in a convex hull of Δ_1-empirical risk minimizers $\{\widehat{g}_i\}_{i=1}^d$. Denote the balls centered at Y to be $\mathcal{B}_1 := \mathcal{B}(Y, \sqrt{r_1})$, $\mathcal{B}_1' := \mathcal{B}(Y, \|\widehat{g} - Y\|_n)$ and $\mathcal{B}_2 := \mathcal{B}(Y, \|\widetilde{f} - Y\|_n)$. The corresponding spheres will be called $\mathcal{S}_1, \mathcal{S}_1', \mathcal{S}_2$. We have $\mathcal{B}_2 \subseteq \mathcal{B}_1$ and $\mathcal{B}_2 \subseteq \mathcal{B}_1'$. Denote by $\mathcal{C}$ the conic hull of $\mathcal{B}_2$ with origin $\widehat{g}$ and define the spherical cap outside the cone $\mathcal{C}$ to be $\mathcal{S} = \mathcal{S}_1' \setminus \mathcal{C}$.

First, $\widetilde{f} \in \mathcal{B}_2$ and it is a contact point of $\mathcal{C}$ and $\mathcal{S}_2$. Indeed, $\widetilde{f}$ is necessarily on a line segment between $\hat{g}$ and a point outside $\mathcal{B}_1$ that does not pass through the interior of $\mathcal{B}_2$ by optimality of $\widetilde{f}$. Let K be the set of all contact points of $\mathcal{C}$ and $\mathcal{S}_2$ – potential locations of $\widetilde{f}$.

Second, for any $h \in \mathcal{F}$, we have $\|h - Y\|_n \geq \sqrt{r_1}$ i.e. any $h \in \mathcal{F}$ is not in the interior of $\mathcal{B}_1$. Furthermore, let $\mathcal{C}'$ be bounded subset cone $\mathcal{C}$ cut at K. Thus $h \in (int\mathcal{C})^c \cap (\mathcal{B}_1)^c$ or $h \in \mathcal{T}$, where $\mathcal{T} := (int\mathcal{C}') \cap (\mathcal{B}_1)^c$.

For any $h \in \mathcal{F}$ consider the two dimensional plane $\mathcal{L}$ that passes through three points $\hat{g}, Y, h$, depicted in Figure 2. Observe that the left-hand side of the desired inequality (9) is constant as $\widetilde{f}$ ranges over K. The maximization of $\|h - f'\|_n^2$ over $f' \in K$ is achieved by $f' \in K \cap \mathcal{L}$. Hence, to prove the desired inequality, we can restrict our attention to the plane $\mathcal{L}$ and f'. Let $h_\perp$ be the projection of h onto the shell $L \cap S_1'$. By the geometry of the cone and triangle inequality we have:

$$\|f' - \widehat{g}\|_n \geq \frac{1}{2}\|\widehat{g} - h_\perp\|_n \geq \frac{1}{2}\left(\|f' - h_\perp\|_n - \|f' - \widehat{g}\|_n\right),$$

and, hence, $\|f' - \widehat{g}\|_n \geq \|f' - h_\perp\|_n/3$. By the Pythagorean theorem,

$$\|h_\perp - Y\|_n^2 - \|f' - Y\|_n^2 = \|\widehat{g} - Y\|_n^2 - \|f' - Y\|_n^2 = \|f' - \widehat{g}\|_n^2 \geq \frac{1}{9}\|f' - h_\perp\|_n^2.$$

We can now extend this claim to h. Indeed, due to the geometry of the projection $h \to h_\perp$ and the fact that $h \in (int\mathcal{C})^c \cap (int\mathcal{B}_1)^c$ or $h \in \mathcal{T}$ there are 2 possibilities:

a) $h \in (\mathcal{B}_1')^c$. Then $\langle h_\perp - Y, h_\perp - h\rangle_n \leq 0$;

b) $h \in \mathcal{B}_1'$. Then, since $h \in (\mathcal{B}_1)^c$, we have

$$\langle h_\perp - Y, h_\perp - h\rangle_n \leq \left(\|h - Y\| + \|h - h_\perp\|\right)\|h - h_\perp\| \leq \|h_\perp - Y\|_n^2 - \|h - Y\|_n^2 \leq \Delta_1.$$

In both cases, the following inequality is true

$$\|h - Y\|_n^2 - \|f' - Y\|_n^2 = \|h_\perp - h\|_n^2 - 2\langle h_\perp - Y, h_\perp - h\rangle_n + \left(\|h_\perp - Y\|_n^2 - \|f' - Y\|_n^2\right)$$

$$\geq \|h_\perp - h\|_n^2 - 2\Delta_1 + \frac{1}{9}\|f' - Y\|_n^2 \geq \frac{1}{18}\|f' - h\|_n^2 - 2\Delta_1.$$

$\square$

Figure 2: The cut surface $\mathcal{L}$

Lemma A.2 (Geometric Inequality for Δ-empirical minimizers). *Let $\hat{g}_1 \ldots \hat{g}_d$ be Δ_1-empirical risk minimizers from the first step of the $Star_d$ procedure, and $\hat{f}$ be the Δ_2-empirical risk minimizer from the second step of the $Star_d$ procedure. Then, for any $h \in \mathcal{F}$ and $c_{A.2} = \frac{1}{36}$ the following inequality holds:*

$$\widehat{\mathbb{E}}(h - Y)^2 - \widehat{\mathbb{E}}(\hat{f} - Y)^2 \geq c_{A.2}\widehat{E}(\hat{f} - h)^2 - 2(1 + c_{A.2})[\Delta_1 + \Delta_2].$$

Proof. Since Lemma A.1 was actually proven for any $f \in K$, let f'' be the closest point to $\hat{f}$ from K. For this f'' the inequality (9) holds. Similarly to Lemma A.1, there are 2 options: either $\hat{f} \in (int\mathcal{C})^c$, or $\hat{f} \in \mathcal{T}$.

a) Let $\hat{f} \in (int\mathcal{C})^c$, then $\langle \hat{f} - f'', f'' - Y \rangle \geq 0$. Since $\hat{f}$ is Δ_2-empirical risk minimizer, we have $\|\hat{f} - f''\|_n^2 + 2\langle \hat{f} - f'', f'' - Y \rangle + \|f'' - Y\|_n^2 = \|\hat{f} - Y\|_n^2 \leq \|f'' - Y\|_n^2 + \Delta_2$. It means, that $\|\hat{f} - f''\|_n^2 \leq \Delta_2$.

b) Let $\hat{f} \in \mathcal{T}$, then by the cosine theorem (as depicted on Figure 2, $\mathcal{L}$ is the two dimensional plane which passes through $\hat{f}, \hat{g}, Y$):

$$\|\hat{f} - f''\|_n^2 = \|f'' - Y\|_n^2 + \|\hat{f} - Y\|_n^2 - 2\|f'' - Y\|_n\|\hat{f} - Y\|_n \cos(\varphi').$$

But $\cos(\varphi') \geq \cos(\varphi) = \frac{\|f'' - Y\|_n}{\|\hat{g} - Y\|_n}$ and $\|\hat{f} - Y\|_n^2 \geq r_1$. Then we have:

$$\|\hat{f} - f''\|_n^2 \leq \Delta_2 + 2\|f'' - Y\|_n^2 \left(1 - \frac{\|\hat{f} - Y\|_n}{\|\hat{g} - Y\|_n}\right)$$

$$\leq \Delta_2 + 2\frac{\|f'' - Y\|_n^2}{\|\hat{g} - Y\|_n} \left(\frac{\|\hat{g} - Y\|_n^2 - \|\hat{f} - Y\|_n^2}{\|\hat{g} - Y\|_n + \|\hat{f} - Y\|_n}\right) \leq \Delta_1 + \Delta_2.$$

Lemma A.1 states:

$$\|h - Y\|_n^2 \geq \|f'' - Y\|_n^2 + c_{A.1}\|f'' - h\|_n^2 - 2\Delta_1.$$

By using the triangle inequality and the convexity of the quadratic function, we can get the following bound

$$\frac{c_{A.1}}{2}\|\hat{f} - h\|_n^2 \leq c_{A.1}\left(\|\hat{f} - f''\|_n^2 + \|f'' - h\|_n^2\right) \leq c_{A.1}[\Delta_2 + \Delta_1] + c_{A.1}\|f'' - h\|_n^2.$$

12

Combining everything together, we get the required result for the constant $c_{A.2} = \frac{c_{A.1}}{2} = \frac{1}{36}$:

$$\widehat{\mathbb{E}}(h-Y)^2 - \widehat{\mathbb{E}}(\widehat{f}-Y)^2 \geq c_{A.2} \cdot \widehat{\mathbb{E}}(\widehat{f}-h)^2 - 2(1+c_{A.2})[\Delta_1 + \Delta_2].$$

$\square$

For convenience, we introduce a Δ-excess risk

$$\mathcal{E}_\Delta(\widehat{g}) := \mathbb{E}(\widehat{g}-Y)^2 - \inf_{f \in \mathcal{F}} \mathbb{E}(f-Y)^2 - 2(1+c_{A.2})[\Delta_1 + \Delta_2],$$

then the following 2 statements are the direct consequences of the corresponding statements from the article [LRS15]. The only difference is that in our case the geometric inequality has terms on the right side with minimization errors Δ_1, Δ_2. Also our definition of the set $\mathcal{H}$ is different, but all that was needed from it was the property that $\widehat{f}$ lies in $\mathcal{H} + f^*$. For brevity, we will not repeat the proofs, but only indicate the numbers of the corresponding results in the titles of the assertions. We will also proceed for statements the proofs for which we slightly modify or use without changes.

Corollary A.3 (Corollary 3). *Conditioned on the data $\{(X_i, Y_i) : 1 \leq i \leq n\}$, we have a deterministic upper bound for the $Star_d$ estimator:*

$$\mathcal{E}_\Delta(\widehat{f}) \leq (\widehat{\mathbb{E}} - \mathbb{E})[2(f^*-Y)(f^*-\widehat{f})] + \mathbb{E}(f^*-\widehat{f})^2 - (1+c_{A.2}) \cdot \widehat{\mathbb{E}}(f^*-\widehat{f})^2.$$

Theorem A.4 (Theorem 4). *The following expectation bound on excess loss of the $Star_d$ estimator holds:*

$$\mathbb{E}\,\mathcal{E}_\Delta(\widehat{f}) \leq (2F' + F(2+c_{A.2})/2) \cdot \mathbb{E}_\sigma \sup_{h \in \mathcal{H}} \left\{ \frac{1}{n} \sum_{i=1}^n 2\sigma_i h(X_i) - c_{A.4} h(X_i)^2 \right\},$$

where $\sigma_1, \ldots \sigma_n$ are independent Rademacher random variables, $c_{A.4} = \min\left\{\frac{c_{A.2}}{4F'}, \frac{c_{A.2}}{4F(2+c_{A.2})}\right\}$, $F = \sup_{f \in \mathcal{F}} |f|_\infty$ and $F' = \sup_{\mathcal{F}} |Y - f|_\infty$ almost surely.

Theorem A.5 (Theorem 7). *Assume the lower isometry bound in Definition 3.2 holds with $\eta_{lib} = c_{A.2}/4$ and some $\delta_{lib} < 1$ and $\mathcal{H}$ is the set defined in 7. Let $\xi_i = Y_i - f^*(X_i)$. Define*

$$A := \sup_{h \in \mathcal{H}} \frac{\mathbb{E}\, h^4}{(\mathbb{E}\, h^2)^2} \text{ and } B := \sup_{X,Y} \mathbb{E}\, \xi^4.$$

Then there exist two absolute constants $c'_{A.5}, c_{\tilde{A}.5} > 0$ (which only depend on $c_{A.2}$), such that

$$\mathbb{P}\left(\mathcal{E}_\Delta(\widehat{f}) > 4u\right) \leq 4\delta_{lib} + 4\,\mathbb{P}\left(\sup_{h \in \mathcal{H}} \frac{1}{n} \sum_{i=1}^n \sigma_i \xi_i h(X_i) - c_{\tilde{A}.5} h(X_i)^2 > u\right)$$

for any

$$u > \frac{32\sqrt{AB}}{c'_{A.5}} \frac{1}{n}$$

as long as $n > \frac{16(1-c'_{A.5})^2 A}{c'^2_{A.5}} \vee n_0(\mathcal{H}, \delta_{lib}, c_{A.2}/4)$.

Lemma A.6 (Lemma 15). *The offset Rademacher complexity for $\mathcal{H}$ is bounded as:*

$$\mathbb{E}_\sigma \sup_{\mathcal{H}} \left\{ \frac{1}{n} \sum_{i=1}^n 2\sigma_i \xi_i h(X_i) - Ch(X_i)^2 \right\} \leq K(C)\varepsilon + M(C) \cdot \frac{\log \mathcal{N}_2(\mathcal{H}, \varepsilon)}{n}$$

and with probability at least $1 - \delta$

$$\sup_{\mathcal{H}} \left\{ \frac{1}{n} \sum_{i=1}^n 2\sigma_i \xi_i h(X_i) - Ch(X_i)^2 \right\} \leq K(C)\varepsilon + M(C) \cdot \frac{\log \mathcal{N}_2(\mathcal{H}, \varepsilon) + \log \frac{1}{\delta}}{n},$$

where

$$K(C) := 2\left(\sqrt{\sum_{i=1}^n \xi^2/n} + C\right), \quad M(C) := \sup_{h \in \mathcal{H} \setminus \{0\}} 4\frac{\sum_{i=1}^n h(X_i)^2 \xi_i^2}{C \sum_{i=1}^n h(X_i)^2}. \tag{10}$$

Proof. Let $N_2(\mathcal{H},\varepsilon)$ be the ε-net of the $\mathcal{H}$ of size at most $\mathcal{N}_2(\mathcal{H},\varepsilon)$ and $v[h]$ be the closest point from this net for function $h \in \mathcal{H}$, i.e. $\|h - v[h]\|_2 \le \varepsilon$. By using the inequality $v[h]_i^2 \le 2\left(h_i^2 + (v[h]_i - h_i)^2\right)$, we can get next upper bound:

$$\left\{ \frac{1}{n} \sum_{i=1}^n 2\sigma_i \xi_i h(\mathbf{X}_i) - Ch(\mathbf{X}_i)^2 \right\}$$

$$\le \left\{ \frac{1}{n} \sum_{i=1}^n 2\sigma_i \xi_i (h(\mathbf{X}_i) - v[h](\mathbf{X}_i)) + C\left(v[h]^2(\mathbf{X}_i)/2 - h^2(\mathbf{X}_i)\right) \right\}$$

$$+ \frac{1}{n} \sup_{v \in N_2(\mathcal{H},\varepsilon)} \left\{ \sum_{i=1}^n 2\sigma_i \xi_i v(\mathbf{X}_i) - \frac{C}{2} v(\mathbf{X}_i)^2 \right\}$$

$$\le 2\varepsilon \left(\sqrt{\sum_{i=1}^n \xi_i^2/n} + C \right) + \frac{1}{n} \sup_{v \in N_2(\mathcal{H},\varepsilon)} \left\{ \sum_{i=1}^n 2\sigma_i \xi_i v(\mathbf{X}_i) - \frac{C}{2} v(\mathbf{X}_i)^2 \right\}.$$

The right summarand is supremum over set of cardinality not more than $\mathcal{N}_2(\mathcal{H},\varepsilon)$. By using Lemma A.11, we acquire the expected estimates. $\square$

We have now obtained, using the offset Rademacher complexity technique, the upper bound on excess risk in terms of the coverage size of the set $\mathcal{H}$. To get the desired result, we need to obtain an upper bound on the size of the cover $\mathcal{H}$ in terms of the size of the cover $\mathcal{F}$.

Lemma A.7. *For any scale $\varepsilon > 0$, the covering number of $\mathcal{F} \subseteq V(L+1) \cdot \mathcal{B}_2$ (where $\mathcal{B}_2$ is a sphere of radius one in space with norm $\|\cdot\|_n$) and that of $\mathcal{H}$ are bounded in the sense:*

$$\log \mathcal{N}_2(\mathcal{F},\varepsilon) \le \log \mathcal{N}_2(\mathcal{H}, \varepsilon) \le (d+2)\left[\log \mathcal{N}_2\left(\mathcal{F}, \frac{\varepsilon}{3(d+1)} \right) + \log \frac{6(d+1)V(L+1)}{\varepsilon} \right].$$

Proof. If we define as $N(\mathcal{F}, \varepsilon)$ the ε-net cardinality no more then $\mathcal{N}(\mathcal{F}, \varepsilon)$, then the following is true: $N(\mathcal{F}_1, \varepsilon_1) + N(\mathcal{F}_1, \varepsilon_2)$ is $(\varepsilon_1 + \varepsilon_2)$-net for $\mathcal{F}_1 + \mathcal{F}_2$. Hence, $\mathcal{N}(\mathcal{F}_1 + \mathcal{F}_2, \varepsilon_1 + \varepsilon_2) \le \mathcal{N}(\mathcal{F}_1, \varepsilon_1) \cdot \mathcal{N}(\mathcal{F}_2, \varepsilon_2)$. With this we can obtain the following upper bound

$$\mathcal{N}_2(\mathcal{H}, \varepsilon) \le \mathcal{N}_2(\mathcal{F} + Hull_d, \varepsilon) \le \mathcal{N}_2\left(\mathcal{F}, \frac{\varepsilon}{3} \right) \cdot \mathcal{N}_2\left(Hull_d, \frac{2\varepsilon}{3} \right).$$

But since $Hull_d$ is the sum of $d+1$ functions from $\mathcal{F}$ with coefficients in $[-1; 1]$, by the inequaility (3), we can cover this with a net of size no more than

$$\left[\mathcal{N}_2\left(\mathcal{F}, \frac{\varepsilon}{3(d+1)} \right) \cdot \frac{6(d+1)V(L+1)}{\varepsilon} \right]^{d+1}.$$

$\square$

Note that to obtain the required orders, we only need coverage with $\varepsilon = 1/n$.

Corollary A.8. *Let $\mathcal{H}$ defined in 7 for $\mathcal{F} = \mathcal{F}(L, \boldsymbol{p}, s)$, then for V defined in 5 holds*

$$\log \mathcal{N}_2\left(\mathcal{H}, \frac{1}{n} \right) \le c_{A.8} d\, s \log\left(VLnd \right),$$

where $c_{A.8}$ is an indepedent constant.

Proof. By lemma A.7 and inequality 4, we have

$$\log \mathcal{N}_2(\mathcal{H}, 1/n) \le (d+2)\left[\log \mathcal{N}_2\left(\mathcal{F}(L, \mathbf{p}, s), \frac{1}{3n(d+1)} \right) + \log 6n(d+1)V(L+1) \right]$$

$$\le (d+2)\left[(s+1)\log\left(2V^2(L+1)(3n(d+1)) \right) + \log\left(6n(d+1)V(L+1) \right) \right].$$

$\square$

We are now fully prepared to prove the two main results.

Theorem A.9. *Let $\widehat{f}$ be a $Star_d$ estimator and $\mathcal{H}$ be the set defined in 7 for $\mathcal{F} = \mathcal{F}(L, \boldsymbol{p}, s)$. The following expectation bound on excess loss holds:*

$$\mathbb{E}\,\mathcal{E}_\Delta(\widehat{f}) \leq 2(F' + V(L+1)) \cdot \left[\frac{K(C)}{n} + M(C) \cdot \frac{c_{A.8}\, d\, s \log\left(V L n\, d\right)}{n} \right],$$

where $K(C)$, $M(C)$ defined in (10) for constants

$$C = \min\left\{ \frac{c_{A.2}}{4F'}, \frac{c_{A.2}}{4V(L+1)(2+c_{A.2})} \right\}, \quad F' = \sup_{\mathcal{F}} |Y - f|_\infty.$$

Proof. By using Theorem A.4 and inequality 3 we have

$$\mathbb{E}\,\mathcal{E}_\Delta(\widehat{f}) \leq (2F' + V(L+1)(2+c_{A.2})/2) \cdot \mathbb{E}_\sigma \sup_{h \in \mathcal{H}} \left\{ \frac{1}{n}\sum_{i=1}^{n} 2\sigma_i h(\mathbf{X}_i) - Ch(\mathbf{X}_i)^2 \right\},$$

where $C = \min\left\{ \frac{c_{A.2}}{4F'}, \frac{c_{A.2}}{4V(L+1)(2+c_{A.2})} \right\}$, $F' = \sup_{\mathcal{F}} |Y - f|_\infty$ almost surely.

By using Lemma A.6 and corollary A.8 we get desired result

$$\mathbb{E}_\sigma \sup_{\mathcal{H}} \left\{ \frac{1}{n}\sum_{i=1}^{n} 2\sigma_i \xi_i h(\mathbf{X}_i) - Ch(\mathbf{X}_i)^2 \right\} \leq \frac{K(C)}{n} + M(C) \cdot \frac{c_{A.8}\, d\, s \log\left(V L n\, d\right)}{n}.$$

$\square$

Theorem A.10. *Let $\widehat{f}$ be a $Star_d$ estimator and let $\mathcal{H}$ be the set defined in 7 for $\mathcal{F} = \mathcal{F}(L, \boldsymbol{p}, s)$. Assume for $\mathcal{H}$ the lower isometry bound in Definition 3.2 holds with $\eta_{lib} = c_{A.2}/4$ and some $\delta_{lib} < 1$. Let $\xi_i = Y_i - f^*(\mathbf{X}_i)$. Define*

$$A := \sup_{h \in \mathcal{H}} \frac{\mathbb{E}\, h^4}{(\mathbb{E}\, h^2)^2} \quad \text{and} \quad B := \sup_{\mathbf{X}, Y} \mathbb{E}\, \xi^4.$$

Then there exist 3 absolute constants $c'_{A.10}, c_{\tilde{A}.10}, c_{A.10} > 0$ (which only depend on $c_{A.2}$), such that

$$\mathbb{P}\left(\mathcal{E}_\Delta(\widehat{f}) > 4D \right) \leq 4(\delta_{lib} + \delta)$$

as long as $n > \frac{16(1-c'_{A.10})^2 A}{c'^2_{A.10}} \vee n_0(\mathcal{H}, \delta_{lib}, c_{A.10}/4)$, where

$$K := \left(\sqrt{\sum_{i=1}^{n} \xi^2/n} + 2c_{\tilde{A}.10} \right), \quad M := \sup_{h \in \mathcal{H} \backslash \{0\}} \frac{\sum_{i=1}^{n} h(\mathbf{X}_i)^2 \xi_i^2}{c_{\tilde{A}.10} \sum_{i=1}^{n} h(\mathbf{X}_i)^2},$$

$$D := \max\left(\frac{K}{n} + M \cdot \frac{c_{A.8}\, d\, s \log\left(V L n\, d\right) + \log\frac{1}{\delta}}{n}, \frac{32\sqrt{AB}}{c'_{A.10}}\frac{1}{n} \right)$$

and $c_{A.8}$ is an independent constant.

Proof. By using Theorem A.5 for any $u > \frac{32\sqrt{AB}}{c'_{A.5}}\frac{1}{n}$ we have

$$\mathbb{P}\left(\mathcal{E}_\Delta(\widehat{f}) > 4u \right) \leq 4\delta_{lib} + 4\,\mathbb{P}\left(\sup_{h \in \mathcal{H}} \frac{1}{n}\sum_{i=1}^{n} \sigma_i \xi_i h(\mathbf{X}_i) - c_{\tilde{A}.5} h(\mathbf{X}_i)^2 > u \right)$$

as long as $n > \frac{16(1-c'_{A.5})^2 A}{c'^2_{A.5}} \vee n_0(\mathcal{H}, \delta_{lib}, c_{A.2}/4)$.

By using Lemmas A.6 and A.8 we have with probability no more than δ for any $C > 0$:

$$\sup_{\mathcal{H}} \left\{ \frac{1}{n}\sum_{i=1}^{n} \sigma_i \xi_i h(\mathbf{X}_i) - \frac{C}{2}h(\mathbf{X}_i)^2 \right\} \geq \frac{K(C)}{2}\varepsilon + \frac{M(C)}{2} \cdot \frac{\log\mathcal{N}_2(\mathcal{H}, \varepsilon) + \log\frac{1}{\delta}}{n},$$

where $K(C)$, $M(C)$ are defined in (10). Combining this inequality for $C = 2c_{\tilde{A}.10} = 2c_{\tilde{A}.5}$ and $c'_{A.10} = c'_{A.5}$, $c_{A.10} = c_{A.2}$ we get the required result. $\square$

15

Lemma A.11 (Lemma 9). *Let $V \subset \mathbb{R}^n$ be a finite set, $|V| = N$. Then, for any $C > 0$:*

$$\mathbb{E}_\sigma \max_{v \in V} \left[\frac{1}{n} \sum_{i=1}^n \sigma_i \xi_i v(\boldsymbol{X}_i) - C v(\boldsymbol{X}_i)^2 \right] \leq M \frac{\log N}{n}.$$

For any $\delta > 0$:

$$\mathbb{P}_\sigma \left(\max_{v \in V} \left[\frac{1}{n} \sum_{i=1}^n \sigma_i \xi_i v(\boldsymbol{X}_i) - C v(\boldsymbol{X}_i)^2 \right] > M \frac{\log N + \log \frac{1}{\delta}}{n} \right) \leq \delta,$$

where

$$M := \sup_{v \in V \setminus \{0\}} \frac{\sum_{i=1}^n v(\boldsymbol{X}_i)^2 \xi_i^2}{2C \sum_{i=1}^n v(\boldsymbol{X}_i)^2}.$$

B   Result Tables

Here we additionally present tables with the results of numerical experiments. Particularly for runs with a small number of $epochs$. It can be observed that the SnapStar algorithm is quite good with a strong budget constraint. The results also include a relatively large run for the FASHION MNIST dataset. At the moment, ClassicStar (new warm-up) takes $10 - 11th$ place in the leaderboard[5] for this dataset. Full versions of the following tables can be found in the repository[6].

Name	d	MSE	MAE	R^2	TRAIN MSE	TIME (sec)
Snap Star (shot warm-up)	5	**10.881±0.575**	**2.229**	**0.869**	1.976	7.8
Snap Star (new warm-up)	5	11.285±0.650	2.283	0.864	2.656	6.6
Snap Ensemble	5	11.862±0.616	2.306	0.858	2.629	6.6
Ensemble	5	12.568±0.878	2.399	0.849	4.220	6.8
Classic Star (no warm-up)	5	11.365±0.410	2.278	0.864	2.978	7.2
Classic Star (new warm-up)	5	12.157±0.822	2.353	0.854	3.320	6.2
Big NN	5	12.068±0.860	2.411	0.855	3.644	4.0
Snap Star (shot warm-up)	4	**11.276±0.582**	**2.269**	**0.865**	2.329	6.2
Snap Star (new warm-up)	4	11.598±0.729	2.292	0.861	2.739	5.0
Snap Ensemble	4	11.819±0.341	2.316	0.858	2.819	5.0
Ensemble	4	12.059±0.614	2.365	0.855	3.732	5.0
Classic Star (no warm-up)	4	11.608±0.722	2.286	0.861	3.198	6.2
Classic Star (new warm-up)	4	11.890±0.966	2.319	0.857	3.093	5.2
Big NN	4	12.556±0.904	2.383	0.849	3.746	4.0

Table 4: BOSTON HOUSE PRICING. Part of results at 30 epochs, $p = 0.1$, $lr = 0.01$

Name	d	MSE	MAE	R2	TRAIN MSE	TIME (sec)
Snap Star (shot warm-up)	5	76.31 ± 0.17	5.97	0.362	70.64	733
Snap Star (new warm-up)	5	76.21 ± 0.10	5.99	0.363	71.34	667
Snap Ensemble	5	76.42 ± 0.11	6.02	0.361	70.03	543
Ensemble	5	76.34 ± 0.07	6.05	0.361	72.05	711
Classic Star (no warm-up)	5	76.57 ± 0.15	6.07	0.36	73.62	783
Classic Star (new warm-up)	5	$\mathbf{76.06 \pm 0.10}$	6.00	0.364	72.59	807
Big NN	5	77.04 ± 0.21	6.02	0.356	75.62	436
Snap Star (shot warm-up)	4	76.30 ± 0.12	5.99	0.362	71.04	632
Snap Star (new warm-up)	4	76.14 ± 0.11	6.01	0.363	71.78	565
Snap Ensemble	4	76.46 ± 0.12	6.02	0.360	70.37	452
Ensemble	4	76.40 ± 0.08	6.05	0.361	72.08	593
Classic Star (no warm-up)	4	76.51 ± 0.04	6.04	0.36	73.76	652
Classic Star (new warm-up)	4	$\mathbf{76.01 \pm 0.10}$	6.01	0.364	72.69	676
Big NN	4	77.06 ± 0.18	6.03	0.355	75.63	375
Snap Star (shot warm-up)	3	76.39 ± 0.32	5.98	0.361	71.62	530
Snap Star (new warm-up)	3	$\mathbf{76.10 \pm 0.07}$	6.00	0.363	72.38	463
Snap Ensemble	3	76.53 ± 0.14	6.02	0.360	70.77	362
Ensemble	3	76.43 ± 0.09	6.05	0.361	72.12	473
Classic Star (no warm-up)	3	76.51 ± 0.12	6.04	0.360	74.00	522
Classic Star (new warm-up)	3	76.16 ± 0.13	6.01	0.363	72.81	546
Big NN	3	76.80 ± 0.23	6.03	0.358	75.60	315

Table 5: MILLIION SONG. Part of results at 10 epochs

---

[5] https://paperswithcode.com/sota/image-classification-on-fashion-mnist
[6] https://github.com/mordiggian174/star-ensembling

Name	d	accuracy	entropy	TIME (sec)
Snap Star (shot warm-up)	3	**0.900±0.002**	**0.284±0.008**	340.333
Snap Star (new warm-up)	3	0.898±0.002	0.285±0.008	313.0
Snap Ensemble	3	0.897±0.003	0.290±0.009	272.667
Ensemble	3	0.887±0.001	0.310±0.005	272.667
Classic Star (no warm-up)	3	0.893±0.002	0.298±0.007	339.667
Classic Star (new warm-up)	3	0.893±0.002	0.297±0.007	285.667
Big NN	3	0.890±0.010	0.299±0.022	214.333
Snap Star (shot warm-up)	2	**0.894±0.007**	**0.294±0.020**	248.667
Snap Star (new warm-up)	2	0.892±0.001	**0.294±0.006**	230.333
Snap Ensemble	2	0.891±0.006	0.302±0.021	203.667
Ensemble	2	0.886±0.004	0.313±0.008	203.0
Classic Star (no warm-up)	2	0.889±0.003	0.304±0.009	249.0
Classic Star (new warm-up)	2	0.889±0.004	0.303±0.008	203.667
Big NN	2	0.892±0.003	0.304±0.007	165.333
Snap Star (shot warm-up)	1	**0.891±0.002**	**0.299±0.006**	159.0
Snap Star (new warm-up)	1	0.885±0.001	0.318±0.008	149.333
Snap Ensemble	1	0.889±0.001	0.304±0.007	136.0
Ensemble	1	0.886±0.005	0.314±0.011	136.333
Classic Star (no warm-up)	1	0.888±0.002	0.311±0.001	158.0
Classic Star (new warm-up)	1	**0.891±0.002**	0.302±0.005	122.333
Big NN	1	0.886±0.002	0.315±0.005	117.333

Table 6: FASHION MNIST. Part of results at 5 epochs, $lr = 0.001$

Name	d	accuracy	entropy	TIME (sec)
Snap Star (shot warm-up)	5	0.898	1.152	2588.0
Snap Star (new warm-up)	5	0.898	1.136	2369.0
Snap Ensemble	5	0.902	0.330	2036.0
Ensemble	5	0.918	0.229	2052.0
Classic Star (no warm-up)	5	0.922	0.229	2589.0
Classic Star (new warm-up)	5	**0.923**	**0.228**	2239.0
Big NN	5	0.910	0.481	1560.0

Table 7: FASHION MNIST. All of results at 25 epochs, $lr = 0.001$